

# CH : - 5 Regression Analysis

## 1. Supervised Learning

### Regression(integer or float)

- Linear
- 1.)simple
- 2.)multiple
- polynomial
- 1.)simple
- 2.)multiple

### Classification(categorical data)

- KNN
- Decision Tree

### equation

- $y = mx + b$  (linear)
- $y = a + bx + cx^1 + dx^2 + ex^3$  (linear)
- $y = a + bx + cx^2 + dx^3$  (polynomial)
- $y = a + bx^1 + cx^2 + dx^3$  (polynomial)

## Simple Linear Regression

### scikit library

- preprocessing

```
In [100]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
df = pd.read_csv("Book1.csv")
df.head()
```

Out[100]:

	cgpa	package
0	6.89	3.26
1	5.12	1.98
2	7.82	3.25
3	7.42	3.67
4	6.94	3.57

```
In [101]: x = df.iloc[:,0:1]
y = df.iloc[:, -1]
# x = df.drop("package",axis=1)
# x = df[['cgpa', 'x1', 'x2']]
# x = df['cgpa']
# y = df['package']
```

```
In [102]: x.shape
```

```
Out[102]: (200, 1)
```

```
In [103]: y.shape
```

```
Out[103]: (200,)
```

- processing

```
In [104]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
x_train.shape
```

```
Out[104]: (160, 1)
```

```
In [105]: help(train_test_split)
```

Help on function `train_test_split` in module `sklearn.model_selection._split`:

`train_test_split(*arrays, **options)`

Split arrays or matrices into random train and test subsets

Quick utility that wraps input validation and

`next(ShuffleSplit().split(X, y))` and application to input data into a single call for splitting (and optionally subsampling) data in a oneliner.

Read more in the :ref:`User Guide <cross\_validation>`.

Parameters

-----

`*arrays` : sequence of indexables with same length / shape[0]

Allowed inputs are lists, numpy arrays, scipy-sparse matrices or pandas dataframes.

`test_size` : float or int, default=None

If float, should be between 0.0 and 1.0 and represent the proportion of the dataset to include in the test split. If int, represents the absolute number of test samples. If None, the value is set to the complement of the train size. If `train_size` is also None, it will

1 be set to 0.25.

`train_size` : float or int, default=None

If float, should be between 0.0 and 1.0 and represent the proportion of the dataset to include in the train split. If int, represents the absolute number of train samples. If None, the value is automatically set to the complement of the test size.

`random_state` : int or RandomState instance, default=None

t. Controls the shuffling applied to the data before applying the split.

Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

`shuffle` : bool, default=True

e Whether or not to shuffle the data before splitting. If `shuffle=False` then stratify must be None.

`stratify` : array-like, default=None

If not None, data is split in a stratified fashion, using this as the class labels.

Returns

-----

`splitting` : list, length=2 \* len(arrays)

List containing train-test split of inputs.

.. versionadded:: 0.16

e If the input is sparse, the output will be a `scipy.sparse.csr_matrix`. Else, output type is the same as the input type.

Examples

-----

```

>>> import numpy as np
>>> from sklearn.model_selection import train_test_split
>>> X, y = np.arange(10).reshape((5, 2)), range(5)
>>> X
array([[0, 1],
       [2, 3],
       [4, 5],
       [6, 7],
       [8, 9]])
>>> list(y)
[0, 1, 2, 3, 4]

>>> X_train, X_test, y_train, y_test = train_test_split(
...     X, y, test_size=0.33, random_state=42)
...
>>> X_train
array([[4, 5],
       [0, 1],
       [6, 7]])
>>> y_train
[2, 0, 3]
>>> X_test
array([[2, 3],
       [8, 9]])
>>> y_test
[1, 4]

>>> train_test_split(y, shuffle=False)
[[0, 1, 2], [3, 4]]

```

```

In [106]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
x_test.shape

```

Out[106]: (40, 1)

```

In [107]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
y_train.shape

```

Out[107]: (160,)

```

In [108]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
y_test.shape

```

Out[108]: (40,)

```

In [109]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
lr=LinearRegression()
model = lr.fit(x_train,y_train) # creating a model or equation
lr.coef_

```

Out[109]: array([0.56004401])

```
In [110]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
lr=LinearRegression()
model = lr.fit(x_train,y_train) # creating a model or equation
lr.intercept_
```


Out[110]: -1.057454060871304

- posting

```
In [111]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
lr=LinearRegression()
model = lr.fit(x_train,y_train) # creating a model or equation
y_pred = lr.predict(x_test)
y_pred
```

Out[111]: array([2.55731678, 3.16105985, 3.90881503, 2.9671974 , 2.3579154 ,  
3.6595633 , 2.97273633, 4.28546208, 3.49893441, 3.32168874,  
3.21644913, 3.49893441, 3.42692836, 1.95357372, 3.16659878,  
2.68471211, 2.92842491, 4.04174928, 3.07243702, 2.9671974 ,  
1.9701905 , 2.82872422, 3.53216798, 3.73156936, 2.13081939,  
2.61270606, 3.37153909, 2.98381418, 2.52962215, 3.43246729,  
3.35492231, 3.60971296, 2.65701747, 3.10567058, 2.9671974 ,  
2.97273633, 2.38561004, 2.96165847, 1.95911264, 3.52662905])

```
In [112]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
lr=LinearRegression()
model = lr.fit(x_train,y_train) # creating a model or equation
y_pred = lr.predict(x_test)
plt.scatter(df['cgpa'],df['package'])
plt.plot(x_test,y_pred)
plt.show()
```



```
In [113]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
import math
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
lr=LinearRegression()
model = lr.fit(x_train,y_train)
y_pred = lr.predict(x_test)
print("MAE",metrics.mean_absolute_error(y_test,y_pred))
print("MSE",metrics.mean_squared_error(y_test,y_pred))
print("RMSE",math.sqrt(metrics.mean_squared_error(y_test,y_pred)))
print("R2score",metrics.r2_score(y_test,y_pred)) # x depend on y
```

```
MAE 0.28331977292553595
MSE 0.13470894990622276
RMSE 0.3670271787023718
R2score 0.6419528555237273
```

```
In [114]: import numpy as np
test_data = np.array([9.3,7.8])
print(test_data)
print(test_data.shape)
```

```
[9.3 7.8]
(2,)
```

```
In [115]: import numpy as np
test_data = test_data.reshape(-1,1)
print(test_data)
print(test_data.shape)
```

```
[[9.3]
 [7.8]]
(2, 1)
```

```
In [116]: lr.predict(test_data)
```

```
Out[116]: array([4.3834217 , 3.48369105])
```

## Multiple Linear Regression

```
In [1]: # Target is selling_price.
import pandas as pd
df = pd.read_csv("car_data.csv")
df.isna().sum()
```

```
Out[1]: Car_Name      0
Year      0
Selling_Price  0
Present_Price  0
Kms_Driven  0
Fuel_Type    0
Seller_Type  0
Transmission 0
Owner        0
dtype: int64
```

```
In [7]: import pandas as pd
df = pd.read_csv("car_data.csv")
df = df.drop("Car_Name",axis=1)
df
```

Out[7]:

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Own
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	
...	...	...	...	...	...	...	...	...
296	2016	9.50	11.60	33988	Diesel	Dealer	Manual	
297	2015	4.00	5.90	60000	Petrol	Dealer	Manual	
298	2009	3.35	11.00	87934	Petrol	Dealer	Manual	
299	2017	11.50	12.50	9000	Diesel	Dealer	Manual	
300	2016	5.30	5.90	5464	Petrol	Dealer	Manual	

301 rows × 8 columns



```
In [8]: df["Age"] = 2025-df["Year"]
df
```

Out[8]:

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Own	Age
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual		11
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual		12
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual		8
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual		14
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual		11
...	...	...	...	...	...	...	...	...	...
296	2016	9.50	11.60	33988	Diesel	Dealer	Manual		9
297	2015	4.00	5.90	60000	Petrol	Dealer	Manual		10
298	2009	3.35	11.00	87934	Petrol	Dealer	Manual		16
299	2017	11.50	12.50	9000	Diesel	Dealer	Manual		8
300	2016	5.30	5.90	5464	Petrol	Dealer	Manual		9

301 rows × 9 columns





```
In [9]: df = df.drop("Year",axis=1)
df
```

Out[9]:

	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	A
0	3.35	5.59	27000	Petrol	Dealer	Manual	0	
1	4.75	9.54	43000	Diesel	Dealer	Manual	0	
2	7.25	9.85	6900	Petrol	Dealer	Manual	0	
3	2.85	4.15	5200	Petrol	Dealer	Manual	0	
4	4.60	6.87	42450	Diesel	Dealer	Manual	0	
...	...	...	...	...	...	...	...	
296	9.50	11.60	33988	Diesel	Dealer	Manual	0	
297	4.00	5.90	60000	Petrol	Dealer	Manual	0	
298	3.35	11.00	87934	Petrol	Dealer	Manual	0	
299	11.50	12.50	9000	Diesel	Dealer	Manual	0	
300	5.30	5.90	5464	Petrol	Dealer	Manual	0	

301 rows × 8 columns



```
In [10]: pd.get_dummies(df)
```

Out[10]:

IG	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type_Dealer	Seller_Type_Individual	Transmission_A
0	0	1	1		0
0	1	0	1		0
0	0	1	1		0
0	0	1	1		0
0	1	0	1		0
...	...	...	...		...
0	1	0	1		0
0	0	1	1		0
0	0	1	1		0
0	1	0	1		0
0	0	1	1		0



```
In [11]: df = pd.get_dummies(df,drop_first=True) # only one column remove from each cat
df
```

Out[11]:

	Selling_Price	Present_Price	Kms_Driven	Owner	Age	Fuel_Type_Diesel	Fuel_Type_Petrol
0	3.35	5.59	27000	0	11	0	1
1	4.75	9.54	43000	0	12	1	0
2	7.25	9.85	6900	0	8	0	1
3	2.85	4.15	5200	0	14	0	1
4	4.60	6.87	42450	0	11	1	0
...	...	...	...	...	...	...	...
296	9.50	11.60	33988	0	9	1	0
297	4.00	5.90	60000	0	10	0	1
298	3.35	11.00	87934	0	16	0	1
299	11.50	12.50	9000	0	8	1	0
300	5.30	5.90	5464	0	9	0	1

301 rows × 9 columns

```
In [12]: # target is selling price
y = df['Selling_Price']
x = df.drop('Selling_Price',axis=1)
x.shape
```

Out[12]: (301, 8)

```
In [13]: y = df['Selling_Price']
x = df.drop('Selling_Price',axis=1)
y.shape
```

Out[13]: (301,)

```
In [30]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_stat
x_train.shape
```

Out[30]: (240, 8)

```
In [16]: x_test.shape
```

Out[16]: (61, 8)

```
In [17]: y_train.shape
```

Out[17]: (240,)

```
In [18]: y_test.shape
```

```
Out[18]: (61,)
```

```
In [31]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
model = lr.fit(x_train,y_train)
lr.coef_
```

```
Out[31]: array([ 4.28802981e-01, -5.92316903e-06, -8.49009727e-01, -4.06411225e-01,
 2.86056285e+00,  7.03103550e-01, -1.07477347e+00, -1.50897975e+00])
```

```
In [32]: print(x.columns)
```

```
Index(['Present_Price', 'Kms_Driven', 'Owner', 'Age', 'Fuel_Type_Diesel',
      'Fuel_Type_Petrol', 'Seller_Type_Individual', 'Transmission_Manual'],
      dtype='object')
```

```
In [33]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
model = lr.fit(x_train,y_train)
lr.intercept_
```

```
Out[33]: 6.850757321843104
```

```
In [52]: from sklearn import metrics
import math
y_pred = lr.predict(x_test)
y_pred
```

```
Out[52]: array([10.57889241,  0.71754255,  4.23613904,  5.17855444,  9.75533583,
 4.20615652,  2.6748089 ,  7.63221623,  0.17032381,  5.13283747,
 6.15727726,  6.44545506,  2.11900126,  7.6644102 ,  1.91710055,
 1.71680788,  2.02672159,  1.85294787,  9.56717323,  4.23076317,
 1.48724915,  9.36189904,  1.46103587,  9.54572517,  0.82954545,
 8.32266051,  1.53698563, -3.16293717,  4.2127377 ,  2.10385587,
 3.42279444,  3.71284702,  5.54538889,  7.6834828 , -1.89757368,
 7.06591847,  8.46027656,  5.77704136,  6.10531922,  6.52919752,
16.08564344,  2.07410694,  1.05462465, -0.44710144,  7.06940454,
 6.73548589,  0.98609923,  7.08961114, 14.16420523,  3.00434598,
 8.2871215 , -0.87862345,  8.8814273 ,  1.15099681,  2.11848057,
-0.82641923,  0.69287067,  9.88583992, -0.45444272, -2.41051663,
10.28627971])
```

```
In [53]: from sklearn import metrics
import math
y_pred = lr.predict(x_test)
print("MAE : ",metrics.mean_absolute_error(y_test,y_pred))
print("MSE : ",metrics.mean_squared_error(y_test,y_pred))
print("RMSE : ",math.sqrt(metrics.mean_squared_error(y_test,y_pred)))
print("R2score : ",metrics.r2_score(y_test,y_pred))
```

```
MAE :  1.2453565634870642
MSE :  2.750689262139753
RMSE :  1.658520202511791
R2score :  0.8502332355855595
```

```
In [54]: # when column value is given
import numpy as np
test_data = np.array([6.7, 10000,0, 7, 1,0, 0, 1])
test_data = test_data.reshape(-1,8)
test_data
```

```
Out[54]: array([[6.7e+00, 1.0e+04, 0.0e+00, 7.0e+00, 1.0e+00, 0.0e+00, 0.0e+00,
                1.0e+00]])
```

```
In [57]: import numpy as np
test_data = np.array([6.7, 10000,0, 7, 1,0, 0, 1])
test_data = test_data.reshape(-1,8)
lr.predict(test_data)
```

```
Out[57]: array([8.17121013])
```

## Advertising csv

```
In [58]: import pandas as pd
df = pd.read_csv("advertising.csv")
df.isna().sum()
```

```
Out[58]: TV          0
Radio          0
Newspaper      0
Sales          0
dtype: int64
```

```
In [66]: y = df["Sales"]
x = df.drop("Sales",axis=1)
```

```
In [67]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
x_train.shape
```

```
Out[67]: (160, 3)
```

```
In [68]: x_test.shape
```

```
Out[68]: (40, 3)
```

```
In [69]: y_train.shape
```

```
Out[69]: (160,)
```

```
In [70]: y_test.shape
```

```
Out[70]: (40,)
```

```
In [74]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
model = lr.fit(x_train,y_train)
lr.coef_
```

```
Out[74]: array([ 0.05413044,  0.1096093 , -0.00071742])
```

```
In [75]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
model = lr.fit(x_train,y_train)
lr.intercept_
```

```
Out[75]: 4.601987623229277
```

```
In [76]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
model = lr.fit(x_train,y_train)
y_pred = lr.predict(x_test)
y_pred
```

```
Out[76]: array([14.06714329, 17.00172214, 20.74018701, 20.14958878, 21.0236248 ,
                16.76572535, 15.76327412,  8.03654982, 17.11011956, 21.34362676,
                11.28608236, 11.06766208, 15.39394793,  8.81462303, 24.55268079,
                17.07362544, 24.88331835, 10.42926552,  6.09130108, 24.92061384,
                9.86831195, 13.35784443, 18.10206165, 10.148075  ,  9.51161825,
                10.5969708 , 15.79896773, 23.52774486, 20.65344938, 18.79356837,
                19.4127532 , 15.18136138, 19.20252368, 12.03755096, 20.51207209,
                15.53065475, 16.36903336, 24.50947859, 14.4166112 , 12.55345811])
```

```
In [77]: from sklearn import metrics
import math
y_pred = lr.predict(x_test)
print("MAE : ",metrics.mean_absolute_error(y_test,y_pred))
print("MSE : ",metrics.mean_squared_error(y_test,y_pred))
print("RMSE : ",math.sqrt(metrics.mean_squared_error(y_test,y_pred)))
print("R2score : ",metrics.r2_score(y_test,y_pred))
```

```
MAE :  1.383308862610568
MSE :  2.886920336419673
RMSE :  1.6990939751584293
R2score :  0.8982801816554089
```

## Simple Ploynomial regression

```
In [78]: import pandas as pd
df = pd.read_csv("temp.csv")
df
```

Out[78]:

	sno	Temperature	Pressure
0	1	0	0.0002
1	2	20	0.0012
2	3	40	0.0060
3	4	60	0.0300
4	5	80	0.0900
5	6	100	0.2700

```
In [80]: df = df.drop("sno",axis=1)
df
```

Out[80]:

	Temperature	Pressure
0	0	0.0002
1	20	0.0012
2	40	0.0060
3	60	0.0300
4	80	0.0900
5	100	0.2700

```
In [81]: y = df["Pressure"]
x = df.drop("Pressure",axis=1)
```

```
In [83]: from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=2)
x_poly = poly.fit_transform(x)
x_poly
```

```
Out[83]: array([[1.0e+00, 0.0e+00, 0.0e+00],
                [1.0e+00, 2.0e+01, 4.0e+02],
                [1.0e+00, 4.0e+01, 1.6e+03],
                [1.0e+00, 6.0e+01, 3.6e+03],
                [1.0e+00, 8.0e+01, 6.4e+03],
                [1.0e+00, 1.0e+02, 1.0e+04]])
```

```
In [84]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x_poly,y,train_size=0.2)
```

```
In [85]: from sklearn.linear_model import LinearRegression  
lr = LinearRegression()  
lr.fit(x_train,y_train)  
lr.coef_
```

Out[85]: array([0., 0., 0.])

```
In [86]: lr.intercept_
```

Out[86]: 0.0002